

Database Management Systems Interview Preparation Guide

Basic Concepts

- **Database:** Organized collection of structured data
- **DBMS:** Software to manage databases; provides interface between data and users
- **Data Independence:** Ability to modify schema without affecting applications
 - **Logical Independence:** Change logical schema without changing applications
 - **Physical Independence:** Change physical storage without changing logical schema
- **Schema:** Database structure definition
 - **Physical Schema:** How data is stored physically
 - **Logical Schema:** Programmers' view of database
 - **View Schema:** End users' view of database

Data Models

- **Relational Model:** Data in tables (relations) with rows and columns
- **Entity-Relationship Model:** Entities, attributes, and relationships
- **Object-Oriented Model:** Data represented as objects
- **Hierarchical Model:** Tree-like structure (parent-child relationships)
- **Network Model:** Records with links between them (graph structure)

Relational Database Concepts

- **Relation:** Table with rows and columns
- **Tuple:** Row in a relation
- **Attribute:** Column in a relation
- **Degree:** Number of attributes in a relation
- **Cardinality:** Number of tuples in a relation
- **Primary Key:** Unique identifier for each tuple
- **Foreign Key:** References primary key of another relation
- **Candidate Key:** Attribute(s) that can uniquely identify a tuple
- **Super Key:** Set of attributes containing a candidate key
- **Composite Key:** Primary key made of multiple attributes

Normalization

- **Purpose:** Minimize redundancy and dependency; prevent anomalies

- **Normal Forms:**

- **1NF:** Atomic values; no repeating groups
- **2NF:** 1NF + no partial dependency of non-key attributes on the key
- **3NF:** 2NF + no transitive dependency of non-key attributes on the key
- **BCNF:** Every determinant is a candidate key
- **4NF:** No multi-valued dependencies
- **5NF:** No join dependencies

SQL (Structured Query Language)

- **DDL (Data Definition Language):** CREATE, ALTER, DROP, TRUNCATE
- **DML (Data Manipulation Language):** SELECT, INSERT, UPDATE, DELETE
- **DCL (Data Control Language):** GRANT, REVOKE
- **TCL (Transaction Control Language):** COMMIT, ROLLBACK, SAVEPOINT

Joins

- **INNER JOIN:** Returns rows when matches exist in both tables
- **LEFT JOIN:** Returns all rows from left table and matching rows from right
- **RIGHT JOIN:** Returns all rows from right table and matching rows from left
- **FULL JOIN:** Returns rows when matches exist in either table
- **CROSS JOIN:** Cartesian product of two tables
- **SELF JOIN:** Joining a table to itself

Indexes

- **Purpose:** Speed up data retrieval
- **Types:**
 - **Primary Index:** On primary key
 - **Secondary Index:** On non-primary fields
 - **Clustered Index:** Determines physical order of data
 - **Non-Clustered Index:** Logical order with pointers to physical data
- **Data Structures:** B-Tree, B+ Tree, Hash Index, Bitmap Index

Transactions

- **ACID Properties:**
 - **Atomicity:** All operations complete or none do
 - **Consistency:** Database moves from one valid state to another

- **Isolation:** Concurrent transactions don't affect each other
- **Durability:** Committed changes are permanent
- **Transaction States:** Active, Partially Committed, Committed, Failed, Aborted
- **Concurrency Control:** Ensures isolation between transactions
- **Locking Mechanisms:** Shared (read) locks, Exclusive (write) locks

Concurrency Control

- **Problems:**
 - **Dirty Read:** Reading uncommitted data
 - **Non-Repeatable Read:** Re-reading returns different values
 - **Phantom Read:** Re-execution of query finds new rows
 - **Lost Update:** Two transactions overwrite each other
- **Solutions:**
 - **Lock-Based:** Two-phase locking (2PL)
 - **Timestamp-Based:** Assign timestamps to transactions
 - **Multiversion Concurrency Control (MVCC):** Multiple versions of data
- **Isolation Levels:**
 - **Read Uncommitted:** Lowest isolation; allows dirty reads
 - **Read Committed:** Prevents dirty reads
 - **Repeatable Read:** Prevents non-repeatable reads
 - **Serializable:** Highest isolation; prevents phantom reads

Deadlocks

- **Definition:** Circular waiting condition where each transaction holds resources needed by others
- **Necessary Conditions** (all four must occur simultaneously):
 - **Mutual Exclusion:** At least one resource must be held in non-sharable mode
 - **Hold and Wait:** Process holding resources can request additional resources
 - **No Preemption:** Resources cannot be forcibly taken from processes
 - **Circular Wait:** Circular chain of processes, each waiting for resource held by next
- **Detection:**
 - **Wait-For Graph:** Directed graph showing which transaction waits for which
 - **Timeout Detection:** Assuming deadlock if transaction waits too long
- **Prevention:**
 - **Requiring all locks at once:** Eliminates hold and wait

- **Resource ordering:** Acquire locks in predefined order to prevent cycles
- **Wait-Die/Wound-Wait:** Based on transaction timestamps
- **Resolution:**
 - **Victim Selection:** Choose transaction to abort (based on age, progress, resources held)
 - **Rollback:** Complete vs partial transaction rollback
 - **Starvation Prevention:** Ensuring no transaction is repeatedly victimized

Database Recovery

- **Types of Failures:**
 - **Transaction Failure:** Logical errors, system errors, deadlocks
 - **System Failure:** Power outage, hardware/software failures
 - **Media Failure:** Disk crash, storage corruption
 - **Network Failure:** Communication breakdown between distributed components
- **Recovery Techniques:**
 - **Log-Based Recovery:** Write-ahead logging (WAL)
 - Undo logging: Records old values before updates
 - Redo logging: Records new values of updates
 - Undo/Redo logging: Records both old and new values
 - **Shadow Paging:** Maintains shadow copy of database
 - **Remote Backup Systems:** Standby database for disaster recovery
- **Recovery Operations:**
 - **Undo:** Reverses uncommitted transactions (rollback)
 - **Redo:** Reapplies committed transactions to restore state
 - **ARIES:** Algorithm for Recovery and Isolation Exploiting Semantics
 - Log sequence numbers (LSNs) for ordering
 - Write-ahead logging protocol
 - Repeating history during redo
 - Logging compensation records during undo
- **Checkpoint Mechanisms:**
 - **Simple Checkpoint:** Pauses all transaction processing
 - **Fuzzy Checkpoint:** Allows transactions during checkpoint
 - **Incremental Checkpoint:** Checkpoints subsets of database
- **Recovery Management:**
 - **Recovery Manager:** Handles database recovery process

- **Buffer Manager:** Manages cached data pages
- **Transaction Manager:** Tracks transaction states
- **High Availability Solutions:**
 - **Replication:** Maintaining multiple copies of data
 - **Failover Systems:** Automatic switching to standby system
 - **Data Mirroring:** Real-time duplication of data
 - **Clustering:** Multiple servers acting as one system

Database Security

- **Authentication:** Verifying user identity
- **Authorization:** Access control through privileges
- **Encryption:** Data encryption at rest and in transit
- **Auditing:** Tracking database access and changes
- **SQL Injection:** Attack using malicious SQL code

Advanced Topics

Distributed Databases

- **Fragmentation:** Horizontal (rows) vs Vertical (columns)
- **Replication:** Copying data to multiple locations
- **Consistency Models:** Strong, Eventual consistency
- **CAP Theorem:** Consistency, Availability, Partition tolerance (pick two)

NoSQL Databases

- **Types:**
 - **Key-Value:** Redis, DynamoDB
 - **Document:** MongoDB, CouchDB
 - **Column-Family:** Cassandra, HBase
 - **Graph:** Neo4j, OrientDB
- **BASE Properties:** Basically Available, Soft state, Eventually consistent
- **When to use:** High volume, flexible schema, horizontal scaling

Data Warehousing

- **OLTP vs OLAP:** Transactional vs Analytical processing
- **Star Schema:** Fact table surrounded by dimension tables

- **Snowflake Schema:** Normalized dimension tables
- **ETL Process:** Extract, Transform, Load
- **Data Mining:** Pattern discovery in large datasets

Common Interview Questions

Basic Concepts

1. **Explain DBMS advantages?** Data sharing, redundancy control, data consistency, security, integrity constraints.
2. **What are database anomalies?** Insertion, deletion, and update anomalies caused by data redundancy.
3. **Entity vs Attribute vs Relationship?** Entity is object, attribute is property, relationship is association between entities.

SQL and Database Design

1. **Difference between DELETE, DROP, and TRUNCATE?** DELETE removes rows (can be rolled back), DROP removes tables (can't be rolled back), TRUNCATE removes all rows quickly (can't be rolled back).
2. **Write SQL query for nth highest salary?** Using subquery:
$$\text{SELECT * FROM Employee e1 WHERE n-1} \\ = (\text{SELECT COUNT(DISTINCT salary) FROM Employee e2 WHERE e2.salary > e1.salary})$$
3. **What is normalization and why use it?** Process to reduce redundancy and dependency; prevents anomalies.

Transactions and Deadlocks

1. **Explain ACID properties?** Atomicity (all or nothing), Consistency (valid state transitions), Isolation (concurrent transactions don't interfere), Durability (persistent after commit).
2. **What is a deadlock and how to prevent it?** Circular wait for resources; prevent by resource ordering or requiring all locks at once.
3. **Phantom read vs dirty read?** Phantom: new rows appear; Dirty: reading uncommitted data.

Indexing and Performance

1. **How does indexing improve performance?** Creates data structure for faster lookups without scanning entire table.
2. **Clustered vs non-clustered index?** Clustered: defines physical order; only one per table. Non-clustered: logical ordering; multiple allowed.
3. **When would you avoid indexes?** Small tables, frequent updates/inserts, columns with low selectivity.

Database Recovery

1. **Explain the importance of write-ahead logging?** Ensures data durability by writing log records before actual data changes.
2. **How does ARIES recovery algorithm work?** Uses logging with LSNs, repeats history during redo, and uses compensation logs during undo.
3. **What's the difference between warm and cold backup?** Warm backup is taken while database is running; cold backup requires shutdown.

Advanced Topics

1. **CAP theorem explanation?** Distributed systems can't simultaneously guarantee Consistency, Availability, and Partition tolerance.
2. **When to use NoSQL over RDBMS?** High volume data, flexible schema, horizontal scaling, eventual consistency acceptable.
3. **OLTP vs OLAP?** OLTP: transactional processing (many short transactions); OLAP: analytical processing (complex queries on historical data).

Optimization Techniques

- **Query Optimization:** Rewriting queries for efficiency
- **Denormalization:** Adding redundancy for performance
- **Partitioning:** Horizontal (sharding) vs Vertical
- **Caching:** Storing frequently accessed data in memory
- **Connection Pooling:** Reusing database connections
- **Proper Indexing:** Creating right indexes on right columns
- ****Avoid SELECT ***:** Specify only needed columns
- **Execution Plan Analysis:** Understanding query performance