

SQL Placement Preparation Notes

1. SQL Basics

- **SQL:** Structured Query Language, used for managing relational databases.
- **Types of SQL commands:**
 - **DDL:** CREATE, ALTER, DROP
 - **DML:** SELECT, INSERT, UPDATE, DELETE
 - **DCL:** GRANT, REVOKE
 - **TCL:** COMMIT, ROLLBACK, SAVEPOINT

2. Data Types

- **INT, VARCHAR(n), CHAR(n), DATE, FLOAT, BOOLEAN**

3. Constraints

- **NOT NULL:** No empty values
- **UNIQUE:** No duplicate values
- **PRIMARY KEY:** NOT NULL + UNIQUE
- **FOREIGN KEY:** References another table
- **CHECK:** Ensures value condition
- **DEFAULT:** Sets default value

4. Joins

- **INNER JOIN:** Common rows only
- **LEFT JOIN:** All from left + matched right
- **RIGHT JOIN:** All from right + matched left
- **FULL OUTER JOIN:** All rows from both
- **SELF JOIN:** Join table with itself

5. Subqueries & CTEs

- **Subquery:** Query inside another query
- **CTE (WITH):** Temporary result set for reuse

6. GROUP BY & Aggregates

- Used with functions: COUNT, SUM, AVG, MAX, MIN
- **HAVING:** Filter after GROUP BY

7. Set Operations

- **UNION:** Combines, removes duplicates
- **UNION ALL:** Combines, keeps duplicates
- **INTERSECT:** Common rows
- **EXCEPT:** Rows in first not in second

8. Indexes

- Improve query speed
- **CREATE INDEX** index_name **ON** table(column);

9. Normalization

- Removes redundancy, increases integrity
- **1NF (First Normal Form):** No repeating groups or arrays. All attributes must contain atomic (indivisible) values.
- **2NF (Second Normal Form):** 1NF + every non-key attribute fully functionally dependent on the entire primary key (eliminates partial dependencies).
- **3NF (Third Normal Form):** 2NF + no transitive dependency (non-key attribute should not depend on another non-key attribute).
- **BCNF (Boyce-Codd Normal Form):** A stronger version of 3NF. Every determinant must be a candidate key.
- **4NF (Fourth Normal Form):** BCNF + no multi-valued dependencies (an attribute should not have multiple independent values).

- **5NF (Fifth Normal Form):** 4NF + no join dependency. Data should be reconstructable from smaller relations without redundancy.

10. Transactions

- **ACID** properties:
 - **Atomicity:** All or nothing
 - **Consistency:** Valid state only
 - **Isolation:** Transactions don't interfere
 - **Durability:** Changes persist after commit

Frequently Asked Interview Questions & Answers

1. Difference between WHERE and HAVING?

- WHERE filters rows before aggregation.
- HAVING filters after aggregation (used with GROUP BY).

2. What is the difference between DELETE, TRUNCATE, and DROP?

- DELETE : Removes rows, can be rolled back.
- TRUNCATE : Removes all rows, faster, cannot be rolled back.
- DROP : Deletes table structure permanently.

3. Explain different types of JOINS with examples.

- INNER JOIN : Matches rows from both tables.
- LEFT JOIN : All left rows + matched right.
- RIGHT JOIN : All right rows + matched left.
- FULL JOIN : All rows from both tables.
- Example: `SELECT * FROM A LEFT JOIN B ON A.id = B.id;`

4. What is a PRIMARY KEY vs FOREIGN KEY?

- PRIMARY KEY : Uniquely identifies each row in a table.
- FOREIGN KEY : Enforces link between two tables.

5. How does indexing improve performance?

- Indexes help locate data quickly, avoiding full table scans.
- Downside: Slower INSERT/UPDATE due to maintenance.

6. Write a query to find the 2nd highest salary.

```
SELECT MAX(salary) FROM employees
WHERE salary < (SELECT MAX(salary) FROM employees);
```

7. What are window functions?

- Functions like `ROW_NUMBER()` , `RANK()` , `DENSE_RANK()` that work over a window of rows.
- Example:

```
SELECT name, salary, RANK() OVER (ORDER BY salary DESC) FROM employees;
```

8. What is normalization and its types?

- Technique to reduce data redundancy and improve data integrity.
- Types:
 - 1NF (First Normal Form): No repeating groups or arrays. All attributes must contain atomic (indivisible) values.
 - 2NF (Second Normal Form): 1NF + every non-key attribute fully functionally dependent on the entire primary key (eliminates partial dependencies).
 - 3NF (Third Normal Form): 2NF + no transitive dependency (non-key attribute should not depend on another non-key attribute).
 - BCNF (Boyce-Codd Normal Form): A stronger version of 3NF. Every determinant must be a candidate key.
 - 4NF (Fourth Normal Form): BCNF + no multi-valued dependencies (an attribute should not have multiple independent values).
 - 5NF (Fifth Normal Form): 4NF + no join dependency. Data should be reconstructable from smaller relations without redundancy.

9. Difference between UNION and UNION ALL?

- `UNION` : Removes duplicates.
- `UNION ALL` : Keeps all duplicates.

10. Explain ACID properties with an example.

- **Atomicity**: All operations succeed or fail.
- **Consistency**: Always valid state (e.g., no negative bank balance).
- **Isolation**: Transactions don't conflict (e.g., booking same seat).
- **Durability**: Changes persist even after power loss.

Practice Queries

-- 1. Second highest salary

```
SELECT MAX(salary) FROM employees WHERE salary < (SELECT MAX(salary) FROM employees);
```

-- 2. Employees with duplicate names

```
SELECT name, COUNT(*) FROM employees GROUP BY name HAVING COUNT(*) > 1;
```

-- 3. Employees in multiple departments

```
SELECT emp_id FROM employee_dept GROUP BY emp_id HAVING COUNT(DISTINCT dept_id) > 1;
```

-- 4. Join employees and departments

```
SELECT e.name, d.dept_name FROM employees e JOIN departments d ON e.dept_id = d.id;
```

SQL Window Functions Guide

What Are Window Functions?

Window functions perform calculations across a set of rows that are related to the current row. Unlike aggregate functions, they don't collapse rows into a single result - each row retains its identity while gaining additional calculated columns.

Basic Syntax:

```
window_function() OVER (  
    [PARTITION BY column1, column2, ...]  
    [ORDER BY column1, column2, ...]  
    [ROWS/RANGE frame_specification]  
)
```

Key Components

- **PARTITION BY:** Divides rows into groups (like GROUP BY but doesn't collapse rows)
- **ORDER BY:** Defines the order for calculations within each partition
- **Frame:** Specifies which rows to include in the calculation (default is all rows in partition)

Sample Data

Let's use this employee table for all examples:

```
CREATE TABLE employees (  
    id INT,  
    name VARCHAR(50),  
    department VARCHAR(50),  
    salary DECIMAL(10,2),  
    hire_date DATE  
);
```

```
INSERT INTO employees VALUES  
(1, 'Alice', 'Engineering', 75000, '2020-01-15'),  
(2, 'Bob', 'Engineering', 80000, '2019-03-20'),  
(3, 'Charlie', 'Engineering', 85000, '2021-06-10'),  
(4, 'Diana', 'Marketing', 65000, '2020-08-12'),  
(5, 'Eve', 'Marketing', 70000, '2019-11-05'),  
(6, 'Frank', 'Marketing', 72000, '2021-02-28'),  
(7, 'Grace', 'Sales', 60000, '2020-04-18'),  
(8, 'Henry', 'Sales', 62000, '2019-07-22'),  
(9, 'Ivy', 'Sales', 68000, '2021-09-14');
```

1. ROW_NUMBER()

Assigns a unique sequential number to each row within a partition.

Example 1: Basic Row Numbering

```
SELECT  
    name,  
    department,  
    salary,  
    ROW_NUMBER() OVER (ORDER BY salary DESC) as overall_rank  
FROM employees;
```

Result:

| name | department | salary | overall_rank |
|---------|-------------|--------|--------------|
| ----- | ----- | ----- | ----- |
| Charlie | Engineering | 85000 | 1 |
| Bob | Engineering | 80000 | 2 |
| Alice | Engineering | 75000 | 3 |
| Frank | Marketing | 72000 | 4 |
| Eve | Marketing | 70000 | 5 |
| Ivy | Sales | 68000 | 6 |
| Diana | Marketing | 65000 | 7 |
| Henry | Sales | 62000 | 8 |
| Grace | Sales | 60000 | 9 |

Example 2: Row Number by Department

```
SELECT
    name,
    department,
    salary,
    ROW_NUMBER() OVER (PARTITION BY department ORDER BY salary DESC) as dept_rank
FROM employees;
```

Result:

| name | department | salary | dept_rank |
|---------|-------------|--------|-----------|
| ----- | ----- | ----- | ----- |
| Charlie | Engineering | 85000 | 1 |
| Bob | Engineering | 80000 | 2 |
| Alice | Engineering | 75000 | 3 |
| Frank | Marketing | 72000 | 1 |
| Eve | Marketing | 70000 | 2 |
| Diana | Marketing | 65000 | 3 |
| Ivy | Sales | 68000 | 1 |
| Henry | Sales | 62000 | 2 |
| Grace | Sales | 60000 | 3 |

2. RANK()

Assigns ranks with gaps for tied values. If two rows have the same value, they get the same rank, and the next rank is skipped.

Example 1: Ranking with Ties

```
-- Let's modify data to show ties
UPDATE employees SET salary = 70000 WHERE name = 'Diana';

SELECT
    name,
    department,
    salary,
    RANK() OVER (ORDER BY salary DESC) as salary_rank
FROM employees;
```

Result:

| name | department | salary | salary_rank |
|---------|-------------|--------|-------------|
| Charlie | Engineering | 85000 | 1 |
| Bob | Engineering | 80000 | 2 |
| Alice | Engineering | 75000 | 3 |
| Frank | Marketing | 72000 | 4 |
| Eve | Marketing | 70000 | 5 |
| Diana | Marketing | 70000 | 5 |
| Ivy | Sales | 68000 | 7 |
| Henry | Sales | 62000 | 8 |
| Grace | Sales | 60000 | 9 |

Example 2: Ranking Within Departments

```
SELECT
    name,
    department,
    salary,
    RANK() OVER (PARTITION BY department ORDER BY salary DESC) as dept_rank
FROM employees;
```

3. DENSE_RANK()

Similar to RANK() but doesn't skip ranks after ties. Consecutive ranks are assigned even when there are ties.

Example 1: Dense Ranking

```
SELECT
    name,
    department,
    salary,
    DENSE_RANK() OVER (ORDER BY salary DESC) as dense_rank
FROM employees;
```

Result:

| name | department | salary | dense_rank |
|---------|-------------|--------|------------|
| Charlie | Engineering | 85000 | 1 |
| Bob | Engineering | 80000 | 2 |
| Alice | Engineering | 75000 | 3 |
| Frank | Marketing | 72000 | 4 |
| Eve | Marketing | 70000 | 5 |
| Diana | Marketing | 70000 | 5 |
| Ivy | Sales | 68000 | 6 |
| Henry | Sales | 62000 | 7 |
| Grace | Sales | 60000 | 8 |

Comparison: ROW_NUMBER vs RANK vs DENSE_RANK

```
SELECT
    name,
    salary,
    ROW_NUMBER() OVER (ORDER BY salary DESC) as row_num,
    RANK() OVER (ORDER BY salary DESC) as rank_val,
    DENSE_RANK() OVER (ORDER BY salary DESC) as dense_rank_val
FROM employees;
```

Result:

| name | salary | row_num | rank_val | dense_rank_val |
|---------|--------|---------|----------|----------------|
| ----- | ----- | ----- | ----- | ----- |
| Charlie | 85000 | 1 | 1 | 1 |
| Bob | 80000 | 2 | 2 | 2 |
| Alice | 75000 | 3 | 3 | 3 |
| Frank | 72000 | 4 | 4 | 4 |
| Eve | 70000 | 5 | 5 | 5 |
| Diana | 70000 | 6 | 5 | 5 |
| Ivy | 68000 | 7 | 7 | 6 |
| Henry | 62000 | 8 | 8 | 7 |
| Grace | 60000 | 9 | 9 | 8 |

4. LAG() and LEAD()

Access values from previous or next rows within the same result set.

Example 1: Compare with Previous Row

```
SELECT
    name,
    department,
    salary,
    LAG(salary, 1) OVER (ORDER BY salary DESC) as prev_salary,
    salary - LAG(salary, 1) OVER (ORDER BY salary DESC) as salary_diff
FROM employees;
```

Result:

| name | department | salary | prev_salary | salary_diff |
|---------|-------------|--------|-------------|-------------|
| ----- | ----- | ----- | ----- | ----- |
| Charlie | Engineering | 85000 | NULL | NULL |
| Bob | Engineering | 80000 | 85000 | -5000 |
| Alice | Engineering | 75000 | 80000 | -5000 |
| Frank | Marketing | 72000 | 75000 | -3000 |
| Eve | Marketing | 70000 | 72000 | -2000 |
| Diana | Marketing | 70000 | 70000 | 0 |
| Ivy | Sales | 68000 | 70000 | -2000 |
| Henry | Sales | 62000 | 68000 | -6000 |
| Grace | Sales | 60000 | 62000 | -2000 |

Example 2: Department-wise Comparison

```
SELECT
    name,
    department,
    salary,
    LAG(salary, 1) OVER (PARTITION BY department ORDER BY salary DESC) as prev_in_dept,
    LEAD(salary, 1) OVER (PARTITION BY department ORDER BY salary DESC) as next_in_dept
FROM employees;
```

5. FIRST_VALUE() and LAST_VALUE()

Get the first or last value in a window frame.

Example 1: Highest and Lowest in Department

```
SELECT
    name,
    department,
    salary,
    FIRST_VALUE(salary) OVER (PARTITION BY department ORDER BY salary DESC) as highest_in_dept,
    LAST_VALUE(salary) OVER (PARTITION BY department ORDER BY salary DESC
                             ROWS BETWEEN UNBOUNDED PRECEDING AND UNBOUNDED FOLLOWING) as lowest_
FROM employees;
```

Result:

| name | department | salary | highest_in_dept | lowest_in_dept |
|---------|-------------|--------|-----------------|----------------|
| Charlie | Engineering | 85000 | 85000 | 75000 |
| Bob | Engineering | 80000 | 85000 | 75000 |
| Alice | Engineering | 75000 | 85000 | 75000 |
| Frank | Marketing | 72000 | 72000 | 65000 |
| Eve | Marketing | 70000 | 72000 | 65000 |
| Diana | Marketing | 65000 | 72000 | 65000 |
| Ivy | Sales | 68000 | 68000 | 60000 |
| Henry | Sales | 62000 | 68000 | 60000 |
| Grace | Sales | 60000 | 68000 | 60000 |

Example 2: First and Last Hire by Department

```
SELECT
    name,
    department,
    hire_date,
    FIRST_VALUE(name) OVER (PARTITION BY department ORDER BY hire_date) as first_hire,
    LAST_VALUE(name) OVER (PARTITION BY department ORDER BY hire_date
                           ROWS BETWEEN UNBOUNDED PRECEDING AND UNBOUNDED FOLLOWING) as last_hire
FROM employees;
```

6. NTILE()

Divides rows into a specified number of approximately equal groups.

Example 1: Salary Quartiles

```
SELECT
    name,
    department,
    salary,
    NTILE(4) OVER (ORDER BY salary DESC) as salary_quartile
FROM employees;
```

Result:

| name | department | salary | salary_quartile |
|---------|-------------|--------|-----------------|
| Charlie | Engineering | 85000 | 1 |
| Bob | Engineering | 80000 | 1 |
| Alice | Engineering | 75000 | 1 |
| Frank | Marketing | 72000 | 2 |
| Eve | Marketing | 70000 | 2 |
| Diana | Marketing | 70000 | 2 |
| Ivy | Sales | 68000 | 3 |
| Henry | Sales | 62000 | 3 |
| Grace | Sales | 60000 | 4 |

Example 2: Department Performance Tiers

```
SELECT
    name,
    department,
    salary,
    NTILE(3) OVER (PARTITION BY department ORDER BY salary DESC) as dept_tier
FROM employees;
```

7. PERCENT_RANK()

Calculates the relative rank as a percentage (0 to 1).

Example 1: Salary Percentile

```
SELECT
    name,
    department,
    salary,
    PERCENT_RANK() OVER (ORDER BY salary DESC) as salary_percentile,
    ROUND(PERCENT_RANK() OVER (ORDER BY salary DESC) * 100, 2) as percentile_pct
FROM employees;
```

Result:

| name | department | salary | salary_percentile | percentile_pct |
|---------|-------------|--------|-------------------|----------------|
| ----- | ----- | ----- | ----- | ----- |
| Charlie | Engineering | 85000 | 0.0 | 0.00 |
| Bob | Engineering | 80000 | 0.125 | 12.50 |
| Alice | Engineering | 75000 | 0.25 | 25.00 |
| Frank | Marketing | 72000 | 0.375 | 37.50 |
| Eve | Marketing | 70000 | 0.5 | 50.00 |
| Diana | Marketing | 70000 | 0.5 | 50.00 |
| Ivy | Sales | 68000 | 0.75 | 75.00 |
| Henry | Sales | 62000 | 0.875 | 87.50 |
| Grace | Sales | 60000 | 1.0 | 100.00 |

8. CUME_DIST()

Calculates the cumulative distribution (percentage of rows with values less than or equal to current row).

Example 1: Cumulative Distribution

```
SELECT
    name,
    department,
    salary,
    CUME_DIST() OVER (ORDER BY salary DESC) as cumulative_dist,
    ROUND(CUME_DIST() OVER (ORDER BY salary DESC) * 100, 2) as cumulative_pct
FROM employees;
```

Result:

| name | department | salary | cumulative_dist | cumulative_pct |
|---------|-------------|--------|-----------------|----------------|
| Charlie | Engineering | 85000 | 0.11 | 11.11 |
| Bob | Engineering | 80000 | 0.22 | 22.22 |
| Alice | Engineering | 75000 | 0.33 | 33.33 |
| Frank | Marketing | 72000 | 0.44 | 44.44 |
| Eve | Marketing | 70000 | 0.67 | 66.67 |
| Diana | Marketing | 70000 | 0.67 | 66.67 |
| Ivy | Sales | 68000 | 0.78 | 77.78 |
| Henry | Sales | 62000 | 0.89 | 88.89 |
| Grace | Sales | 60000 | 1.0 | 100.00 |

9. Aggregate Window Functions

Use SUM(), COUNT(), AVG(), MIN(), MAX() with OVER clause for running calculations.

Example 1: Running Totals

```
SELECT
    name,
    department,
    salary,
    SUM(salary) OVER (ORDER BY hire_date ROWS BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW) as running_total,
    AVG(salary) OVER (ORDER BY hire_date ROWS BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW) as running_avg
FROM employees
ORDER BY hire_date;
```

Example 2: Moving Averages

```
SELECT
    name,
    department,
    salary,
    hire_date,
    AVG(salary) OVER (ORDER BY hire_date ROWS BETWEEN 2 PRECEDING AND CURRENT ROW) as moving_avg,
    COUNT(*) OVER (ORDER BY hire_date ROWS BETWEEN 2 PRECEDING AND CURRENT ROW) as window_size
FROM employees
ORDER BY hire_date;
```

10. NTH_VALUE()

Returns the nth value in the window frame.

Example 1: Second Highest Salary

```
SELECT
    name,
    department,
    salary,
    NTH_VALUE(salary, 2) OVER (PARTITION BY department ORDER BY salary DESC
                               ROWS BETWEEN UNBOUNDED PRECEDING AND UNBOUNDED FOLLOWING) as second_highest_salary
FROM employees;
```


Example 2: Middle Value

```
SELECT
    name,
    department,
    salary,
    NTH_VALUE(salary, 2) OVER (PARTITION BY department ORDER BY salary
                               ROWS BETWEEN UNBOUNDED PRECEDING AND UNBOUNDED FOLLOWING) as middle_value
FROM employees;
```

Practical Use Cases

1. Top N per Group

Find top 2 earners in each department:

```
SELECT name, department, salary
FROM (
    SELECT
        name,
        department,
        salary,
        ROW_NUMBER() OVER (PARTITION BY department ORDER BY salary DESC) as rn
    FROM employees
) ranked
WHERE rn <= 2;
```

2. Percentage Calculations

Calculate salary as percentage of department total:

```
SELECT
    name,
    department,
    salary,
    ROUND(salary * 100.0 / SUM(salary) OVER (PARTITION BY department), 2) as pct_of_dept_total
FROM employees;
```

3. Running Totals and Cumulative Analysis

Calculate cumulative salary by department:

```
SELECT
    name,
    department,
    salary,
    SUM(salary) OVER (PARTITION BY department ORDER BY hire_date
                      ROWS BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW) as running_total
FROM employees;
```

4. Year-over-Year Growth Analysis

Track salary changes over time:

```
SELECT
    name,
    department,
    salary,
    hire_date,
    LAG(salary, 1) OVER (PARTITION BY name ORDER BY hire_date) as prev_salary,
    salary - LAG(salary, 1) OVER (PARTITION BY name ORDER BY hire_date) as salary_increase,
    ROUND((((salary - LAG(salary, 1) OVER (PARTITION BY name ORDER BY hire_date)) /
            LAG(salary, 1) OVER (PARTITION BY name ORDER BY hire_date)) * 100, 2) as growth_rate
FROM employee_history;
```

5. Moving Averages for Trend Analysis

Calculate 3-month moving average of sales:

```
SELECT
    month,
    sales_amount,
    AVG(sales_amount) OVER (ORDER BY month ROWS BETWEEN 2 PRECEDING AND CURRENT ROW) as moving_avg,
    sales_amount - AVG(sales_amount) OVER (ORDER BY month ROWS BETWEEN 2 PRECEDING AND CURRENT ROW) as deviation
FROM monthly_sales;
```

6. Gap Analysis - Finding Missing Values

Identify gaps in sequential data:

```
SELECT
    id,
    LAG(id) OVER (ORDER BY id) as prev_id,
    id - LAG(id) OVER (ORDER BY id) as gap
FROM orders
WHERE id - LAG(id) OVER (ORDER BY id) > 1;
```

7. Quartile Analysis for Performance Review

Categorize employees into performance quartiles:

```
SELECT
    name,
    department,
    performance_score,
    NTILE(4) OVER (ORDER BY performance_score DESC) as performance_quartile,
    CASE
        WHEN NTILE(4) OVER (ORDER BY performance_score DESC) = 1 THEN 'Top Performer'
        WHEN NTILE(4) OVER (ORDER BY performance_score DESC) = 2 THEN 'High Performer'
        WHEN NTILE(4) OVER (ORDER BY performance_score DESC) = 3 THEN 'Average Performer'
        ELSE 'Needs Improvement'
    END as performance_category
FROM employees;
```

8. First and Last Transaction Analysis

Track customer behavior patterns:

```

SELECT
    customer_id,
    transaction_date,
    amount,
    FIRST_VALUE(amount) OVER (PARTITION BY customer_id ORDER BY transaction_date) as first_purcl
    LAST_VALUE(amount) OVER (PARTITION BY customer_id ORDER BY transaction_date
                                ROWS BETWEEN UNBOUNDED PRECEDING AND UNBOUNDED FOLLOWING) as last_pi
    ROW_NUMBER() OVER (PARTITION BY customer_id ORDER BY transaction_date) as transaction_sequer
FROM transactions;

```

9. Duplicate Detection and Ranking

Find and rank duplicate records:

```

SELECT
    email,
    name,
    registration_date,
    ROW_NUMBER() OVER (PARTITION BY email ORDER BY registration_date) as duplicate_rank
FROM users
WHERE ROW_NUMBER() OVER (PARTITION BY email ORDER BY registration_date) > 1;

```

10. Median Calculation

Calculate median salary by department:

```
SELECT DISTINCT
    department,
    PERCENTILE_CONT(0.5) WITHIN GROUP (ORDER BY salary) OVER (PARTITION BY department) as median_salary
FROM employees;
```

-- Alternative using NTILE for approximate median

```
SELECT
    department,
    salary,
    NTILE(2) OVER (PARTITION BY department ORDER BY salary) as half,
    CASE
        WHEN NTILE(2) OVER (PARTITION BY department ORDER BY salary) = 1
        THEN 'Lower Half'
        ELSE 'Upper Half'
    END as salary_group
FROM employees;
```

11. Cohort Analysis

Analyze user retention by signup month:

```
SELECT
    signup_month,
    activity_month,
    COUNT(DISTINCT user_id) as active_users,
    FIRST_VALUE(COUNT(DISTINCT user_id)) OVER (PARTITION BY signup_month ORDER BY activity_month) as first_signup_users,
    ROUND(COUNT(DISTINCT user_id) * 100.0 /
        FIRST_VALUE(COUNT(DISTINCT user_id)) OVER (PARTITION BY signup_month ORDER BY activity_month)) as retention_rate
FROM user_activity
GROUP BY signup_month, activity_month;
```

12. ABC Analysis (Pareto Analysis)

Classify products by revenue contribution:

```

SELECT
    product_id,
    revenue,
    SUM(revenue) OVER (ORDER BY revenue DESC ROWS BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW) :
    SUM(revenue) OVER () as total_revenue,
    ROUND(SUM(revenue) OVER (ORDER BY revenue DESC ROWS BETWEEN UNBOUNDED PRECEDING AND CURRENT
        SUM(revenue) OVER (), 2) as cumulative_percentage,
CASE
    WHEN ROUND(SUM(revenue) OVER (ORDER BY revenue DESC ROWS BETWEEN UNBOUNDED PRECEDING AND
        SUM(revenue) OVER (), 2) <= 80 THEN 'A'
    WHEN ROUND(SUM(revenue) OVER (ORDER BY revenue DESC ROWS BETWEEN UNBOUNDED PRECEDING AND
        SUM(revenue) OVER (), 2) <= 95 THEN 'B'
    ELSE 'C'
END as abc_category
FROM product_sales;

```

13. Time Series Analysis - Peak Detection

Identify local maxima in time series data:

```

SELECT
    date,
    value,
    LAG(value) OVER (ORDER BY date) as prev_value,
    LEAD(value) OVER (ORDER BY date) as next_value,
CASE
    WHEN value > LAG(value) OVER (ORDER BY date) AND
        value > LEAD(value) OVER (ORDER BY date) THEN 'Peak'
    WHEN value < LAG(value) OVER (ORDER BY date) AND
        value < LEAD(value) OVER (ORDER BY date) THEN 'Valley'
    ELSE 'Normal'
END as trend_point
FROM time_series_data;

```

14. Customer Lifetime Value Calculation

Calculate CLV using window functions:

```

SELECT
    customer_id,
    order_date,
    order_value,
    SUM(order_value) OVER (PARTITION BY customer_id ORDER BY order_date
                           ROWS BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW) as lifetime_value,
    COUNT(*) OVER (PARTITION BY customer_id ORDER BY order_date
                   ROWS BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW) as order_count,
    AVG(order_value) OVER (PARTITION BY customer_id ORDER BY order_date
                           ROWS BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW) as avg_order_value
FROM orders;

```

15. Session Analysis

Analyze user session patterns:

```

SELECT
    user_id,
    page_view_time,
    LAG(page_view_time) OVER (PARTITION BY user_id ORDER BY page_view_time) as prev_page_time,
    page_view_time - LAG(page_view_time) OVER (PARTITION BY user_id ORDER BY page_view_time) as
CASE
    WHEN page_view_time - LAG(page_view_time) OVER (PARTITION BY user_id ORDER BY page_view_
    THEN 'New Session'
    ELSE 'Same Session'
END as session_indicator
FROM page_views;

```

Key Differences Summary

| Function | Behavior with Ties | Use Case | Example Result |
|--------------|-----------------------|--------------------------------|------------------|
| ROW_NUMBER() | Always unique numbers | Pagination, unique identifiers | 1, 2, 3, 4, 5, 6 |
| RANK() | Same rank, skips next | Traditional ranking | 1, 2, 2, 4, 5, 6 |
| DENSE_RANK() | Same rank, no gaps | Continuous ranking | 1, 2, 2, 3, 4, 5 |

| Function | Behavior with Ties | Use Case | Example Result |
|----------------|-------------------------|-----------------------|----------------------------------|
| PERCENT_RANK() | Percentage rank (0-1) | Percentile analysis | 0.0, 0.2, 0.4, 0.6, 0.8, 1.0 |
| CUME_DIST() | Cumulative distribution | Distribution analysis | 0.17, 0.33, 0.5, 0.67, 0.83, 1.0 |
| NTILE(n) | Divides into n groups | Quartiles/deciles | 1, 1, 2, 2, 3, 3 |
| LAG()/LEAD() | Previous/next values | Trend analysis | Previous or next row value |
| FIRST_VALUE() | First value in window | Baseline comparison | First value in partition |
| LAST_VALUE() | Last value in window | Final comparison | Last value in partition |
| NTH_VALUE() | Nth value in window | Specific position | Value at nth position |

Window Frame Specifications

Frame Types

- **ROWS:** Physical number of rows
- **RANGE:** Logical range based on values

Common Frame Patterns

```
-- Current row only (default for ranking functions)
ROWS BETWEEN CURRENT ROW AND CURRENT ROW

-- All preceding rows and current row
ROWS BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW

-- All rows in partition
ROWS BETWEEN UNBOUNDED PRECEDING AND UNBOUNDED FOLLOWING

-- Moving window (3 rows)
ROWS BETWEEN 2 PRECEDING AND CURRENT ROW

-- Centered window
ROWS BETWEEN 1 PRECEDING AND 1 FOLLOWING
```

Frame Examples

```
SELECT
    name,
    salary,
    -- Running total
    SUM(salary) OVER (ORDER BY salary ROWS BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW) as runn:
    -- 3-row moving average
    AVG(salary) OVER (ORDER BY salary ROWS BETWEEN 2 PRECEDING AND CURRENT ROW) as moving_avg_3,
    -- Centered average
    AVG(salary) OVER (ORDER BY salary ROWS BETWEEN 1 PRECEDING AND 1 FOLLOWING) as centered_avg
FROM employees;
```

Best Practices

1. Choose the right function for your use case:

- Use **ROW_NUMBER()** when you need unique identifiers or for pagination
- Use **RANK()** for traditional ranking where gaps after ties are acceptable
- Use **DENSE_RANK()** when you want continuous ranking without gaps
- Use **NTILE()** for creating equal-sized groups or percentiles
- Use **LAG()/LEAD()** for comparing with adjacent rows
- Use **FIRST_VALUE()/LAST_VALUE()** for baseline comparisons

2. **Always consider the ORDER BY clause** - it determines the ranking logic and is crucial for meaningful results
3. **Use PARTITION BY strategically** to create separate ranking groups and avoid unexpected results
4. **Be explicit with frame specifications** when using aggregate window functions to ensure predictable results
5. **Handle NULL values appropriately:**

```
-- Use COALESCE or NULLIF when needed  
LAG(salary, 1, 0) OVER (ORDER BY hire_date) as prev_salary_or_zero
```

6. **Combine window functions for complex analysis:**

```
SELECT  
    name,  
    salary,  
    DENSE_RANK() OVER (ORDER BY salary DESC) as salary_rank,  
    NTILE(4) OVER (ORDER BY salary DESC) as salary_quartile,  
    PERCENT_RANK() OVER (ORDER BY salary DESC) as salary_percentile  
FROM employees;
```

7. **Use CTEs for complex window function queries** to improve readability:

```
WITH ranked_employees AS (  
    SELECT  
        name,  
        department,  
        salary,  
        ROW_NUMBER() OVER (PARTITION BY department ORDER BY salary DESC) as dept_rank  
    FROM employees  
)  
SELECT * FROM ranked_employees WHERE dept_rank <= 3;
```

8. **Consider performance implications:**

- Window functions can be resource-intensive on large datasets
- Consider using appropriate indexes on PARTITION BY and ORDER BY columns
- Use LIMIT with caution as it's applied after window functions

Performance Tips

1. **Index strategy:** Create indexes on columns used in PARTITION BY and ORDER BY clauses

```
CREATE INDEX idx_emp_dept_salary ON employees(department, salary DESC);
```

2. **Limit result sets early:** Use WHERE clauses before window functions when possible

```
SELECT
    name,
    ROW_NUMBER() OVER (ORDER BY salary DESC) as rank
FROM employees
WHERE department = 'Engineering' -- Filter first
AND salary > 50000;
```

3. **Use appropriate frame specifications:** Avoid UNBOUNDED FOLLOWING when not necessary

```
-- More efficient
SUM(salary) OVER (ORDER BY date ROWS BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW)

-- Less efficient if you don't need the full window
SUM(salary) OVER (ORDER BY date ROWS BETWEEN UNBOUNDED PRECEDING AND UNBOUNDED FOLLOWING)
```

4. **Consider materialized views** for frequently used window function results
5. **Use EXPLAIN ANALYZE** to understand query execution plans and optimize accordingly

Common Pitfalls to Avoid

1. **Forgetting frame specifications with aggregate functions:**

```
-- This might not give expected results
LAST_VALUE(salary) OVER (PARTITION BY department ORDER BY salary DESC)

-- Better approach
LAST_VALUE(salary) OVER (PARTITION BY department ORDER BY salary DESC
                        ROWS BETWEEN UNBOUNDED PRECEDING AND UNBOUNDED FOLLOWING)
```

2. **Misunderstanding NULL handling:**

```
-- NULLs are typically ordered last in ascending order
ROW_NUMBER() OVER (ORDER BY salary NULLS FIRST) -- Be explicit
```

3. **Using wrong function for the task:**

```
-- Don't use RANK() if you need unique identifiers
-- Use ROW_NUMBER() instead
```

4. Not considering ties in ranking:

- Be aware that RANK() and DENSE_RANK() handle ties differently
- Choose based on your business requirements

5. Inefficient frame specifications:

- Avoid this pattern for large datasets

AVG(salary) OVER (ORDER BY id ROWS BETWEEN UNBOUNDED PRECEDING AND UNBOUNDED FOLLOWING)

- Consider if you really need the full window

AVG(salary) OVER (PARTITION BY department) -- This might be sufficient